

Autonomous Research Agent

Plan, gather, reflect, and synthesize research reports.

HISTORY

- Fable 5 and ops 4.8 comparison ...
- Fable 5 and ops 4.8 comparison ...
- Fable 5 and ops 4.8 comparison ...
- Fable 5 model News of claude an...
- Fable 5 model News of claude an...
- Recent breakthroughs in solid-st...

Provider

openrouter

Model

free

NVIDIA: Nemotron 3 Nano Omni (free) · Free

API key (optional — overrides the server's .env key)

Load models

Session-only: held in your browser, sent to your local server per request, never stored. Local use only.

Tavily API key (enables the Tavily source)

NewsAPI key (enables the News source)

Sources (leave all unchecked to let the agent auto-pick)

☒ Tavily (add key) ☐ Wikipedia ☐ News (add key) ☐ Web

Research question

Fable 5 and ops 4.8 comparison of both claude model

Research

PROGRESS

- plan — sources: tavily, wikipedia, web
- gather — gathered 22 results (iteration 1)
- process — kept 7 results
- reflect — enough=False
- gather — gathered 44 results (iteration 2)
- process — kept 4 results
- reflect — enough=False
- synthesize — drafted report (4337 chars)
- persist — saved run 6
- Done.

REPORT [Export MD](#) [Export PDF](#)

Key Points

Claude **Fable 5** introduces a new **Mythos tier** that sits above **Opus 4.8** in capability. Fable 5 costs **exactly double** the price of Opus 4.8 (**\$10/\$50 per million tokens**). On the **SWE-Bench Pro** benchmark, Fable 5 scores **80.3%** versus Opus 4.8's **69.2%**, widening the performance gap on hard tasks. The model is **slower** and **more expensive**, but provides higher capability and safety considerations. A decision guide evaluates **capability, price, speed, safety**, and offers a **routing strategy** that can use both models together.

Fable 5 was **banned by the US government** for 19 days due to a jailbreak discovered by Amazon researchers, then **reinstated** on July 1. Testing in a side-by-side YouTube video showed **direct, prompt-matched comparisons** with Opus 4.8.

The Mythos tier (Fable 5) is described as a **full capability class above** Opus 4.8, while **Sonnet 4.6** is recommended for daily tasks, Fable 5 for complex jobs, and Opus 4.8 as a fallback.

Important Findings

Performance advantage: Fable 5 outperforms Opus 4.8 on the SWE-Bench Pro benchmark by **~11 percentage points**, indicating a substantial leap in reasoning and coding ability, especially on difficult tasks.

Cost-performance trade-off: The **double price** of Fable 5 may be justified only when the higher accuracy and broader capability are critical; for routine or budget-constrained use, Opus 4.8 remains attractive.

Speed penalty: Users should expect **lower throughput** with Fable 5, which could affect real-time applications or high-volume workloads.

Regulatory risk: The brief ban highlights potential **regulatory scrutiny** of advanced models; organizations must monitor policy changes and ensure compliance before deployment.

Use-case guidance:

Sonnet 4.6 is optimal for **everyday tasks** where speed and cost matter.

Fable 5 is best suited for **complex, high-stakes jobs** (e.g., advanced reasoning, large-scale coding, research) where its superior benchmark scores justify the expense and slower latency.

Opus 4.8 serves as a **reliable fallback** for tasks that do not require the extra capability of Fable 5 but still need strong performance.

Routing strategy: Leveraging both models—using Fable 5 for demanding tasks and falling back to Opus 4.8 for less demanding ones—can balance **performance, cost, and latency** effectively.

References / Sources

[0] tavily – *Claude Fable 5 vs Opus 4.8: Upgrade Guide* - LushBinary – <https://lushbinary.com/blog/claude-fable-5-vs-opus-4-8-upgrade-guide>

[1] tavily – *I Tested The Banned AI Model: Fable 5 vs Opus 4.8* - YouTube – <https://www.youtube.com/watch?v=1TNSJ1gLzFk>

[2] tavily – *Claude Fable 5 vs Opus 4.8: The Mythos Hype Meets Reality* - Tessl – <https://tessl.io/blog/claude-fable-5-vs-opus-48-the-mythos-hype-meets-reality>

[3] tavily – *Claude Fable 5 vs Opus 4.8 vs Sonnet 4.6 (2026)* – <https://vanbeaumont.nl/en/blog/claude-fable-5-opus-4-8-sonnet-4-6-which-model>

Actionable Insights

Assess task complexity: Deploy **Fable 5** only for tasks that demand the highest reasoning accuracy (e.g., complex code generation, research synthesis). For routine or cost-sensitive workloads, stick with **Opus 4.8** or **Sonnet 4.6**.

Budget planning: Factor in the **double token cost** of Fable 5; run cost-benefit analyses to determine whether the performance gain translates to measurable ROI for your use case.

Performance monitoring: Because Fable 5 is **slower**, schedule it for off-peak periods or batch processing where latency is less critical.

Regulatory compliance: Keep abreast of **government regulations** concerning advanced AI models; the temporary ban of Fable 5 suggests that compliance checks are essential before large-scale deployment.

Hybrid routing: Implement a **routing layer** that selects Fable 5 for high-difficulty prompts and falls back to Opus 4.8 (or Sonnet 4.6) for simpler requests, optimizing both **quality and cost**.

Safety and alignment: Leverage the **enhanced safety mechanisms** noted in the decision guide when handling sensitive or high-risk content, as Fable 5's higher capability may also bring increased responsibility.